

KaggleDSF

September 25, 2018

1 New York City Taxi Fare Prediction

The aim is to predict the taxi fare for the customer taking cab service in New York City

Data Import and Exploration

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import gc
import os
import seaborn as sb
%matplotlib inline
```

```
In [2]: print(os.listdir('./'))
```

```
['.config', 'Anaconda3-5.2.0-Linux-x86_64.sh', 'train.csv', 'anaconda3', '.jupyter', 'output_37_
```

```
In [3]: train_df = pd.read_csv('./train.csv')
train_df.dtypes
```

```
Out[3]: key                object
fare_amount              float64
pickup_datetime          object
pickup_longitude         float64
pickup_latitude          float64
dropoff_longitude        float64
dropoff_latitude         float64
passenger_count          int64
dtype: object
```

```
In [4]: print(train_df.shape[0])
```

```
55423856
```

```
In [5]: #Drop any values that are NaN
        train_df = train_df.dropna(how = 'any', axis = 'rows')
        print('New size: %d' % len(train_df))
```

New size: 55423480

```
In [6]: #Drop all the rows that have value = 0 in them
        train_df = train_df[(train_df != 0).all(1)]
        print('New size: %d' % len(train_df))
```

New size: 54127059

```
In [7]: #Take only those Fares of passengers that have a postive value
        train_df = train_df[train_df.fare_amount > 0]
        print('New size: %d' % len(train_df))
```

New size: 54124853

```
In [8]: #Using the general fare_amount from the given data set, it seems the fare of more than 1
        train_df = train_df[train_df.fare_amount < 175]
        print('New size: %d' % len(train_df))
```

New size: 54122340

```
In [9]: #We are taking percentile of the data in order to identify the outliers and make sure to
        train_df['pickup_longitude'].quantile([0, 0.01, 0.05, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7,
```

```
Out[9]: 0.00    -3442.059565
        0.01     -74.014384
        0.05     -74.006905
        0.10     -74.002911
        0.20     -73.994792
        0.30     -73.990328
        0.40     -73.986191
        0.50     -73.982087
        0.60     -73.977850
        0.70     -73.972158
        0.80     -73.963565
        0.90     -73.953194
        0.95     -73.934200
        0.99     -73.785999
        1.00     3457.625683
        Name: pickup_longitude, dtype: float64
```

- From the above values we can see that the majority of data lies between 1%tile and 99%tile. So we are now going to filter the data according to their percentile ranges

```

In [10]: print('Size of training data: %d' % len(train_df) )

train_df = train_df[(train_df['pickup_longitude'] <= train_df['pickup_longitude'].quantil
train_df = train_df[(train_df['pickup_longitude'] >= train_df['pickup_longitude'].quantil

train_df = train_df[(train_df['pickup_latitude'] <= train_df['pickup_latitude'].quantil
train_df = train_df[(train_df['pickup_latitude'] >= train_df['pickup_latitude'].quantil

train_df = train_df[(train_df['dropoff_longitude'] <= train_df['dropoff_longitude'].qua
train_df = train_df[(train_df['dropoff_longitude'] >= train_df['dropoff_longitude'].qua

train_df = train_df[(train_df['dropoff_latitude'] <= train_df['dropoff_latitude'].quant
train_df = train_df[(train_df['dropoff_latitude'] >= train_df['dropoff_latitude'].quant

print('New size: %d' % len(train_df))

```

Size of training data: 54122340
New size: 52852731

We tried taking different values of the percentile, and taking the data between 0.5%tile and 99.9%tile gives us the best data set for training.

```

In [11]: eu_cal = (train_df['dropoff_latitude'] - train_df['pickup_latitude']) **2 + (train_df[

eu_dist = np.sqrt(eu_cal)

```

1.0.1 Pearson Correlation between Euclidian Distance and Fare Amount

```

In [12]: eu_dist.corr(train_df['fare_amount'])

```

Out[12]: 0.8806340289981661

```

In [13]: #Get the given pickup time in a new format of date of Y-M-D:H-M-S
train_df['pickup_datetime'] = train_df['pickup_datetime'].str.replace(" UTC", "")
#replace the given date time in a new format
train_df['pickup_datetime'] = pd.to_datetime(train_df['pickup_datetime'], format='%Y-%m

```

```

In [14]: journey_time = (train_df['pickup_datetime'].dt.hour)*60 + train_df['pickup_datetime'].d

```

1.0.2 Pearson Correlation between Euclidian Distance and the Time

```

In [15]: eu_dist.corr(journey_time)

```

Out[15]: -0.032817644627646574

1.0.3 Pearson Correlation between Fare Amount and the Time

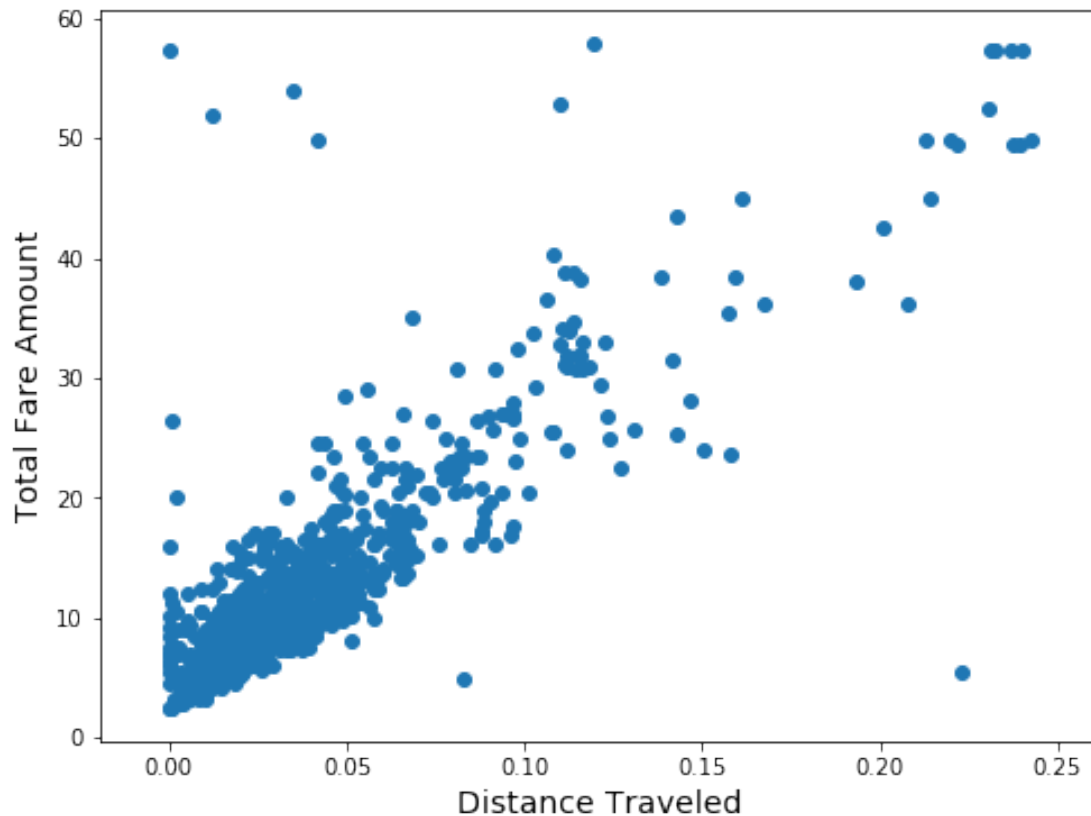
```
In [16]: journey_time.corr(train_df['fare_amount'])
```

```
Out[16]: -0.01736950684640317
```

```
In [17]: train_df['distance'] = eu_dist  
         train_df['journey_time'] = journey_time
```

1.0.4 Plot between Distance and Total Fare Amount


```
In [18]: plt.figure(figsize = (8,6))  
         plt.xlabel('Distance Traveled', fontsize = 14)  
         plt.ylabel('Total Fare Amount', fontsize = 14)  
         plt.scatter(train_df[:1200].distance, train_df[:1200].fare_amount)  
         plt.show()
```



- The plot between Distance traveled and the Total Fare Amount generates a linear relationship

1.0.5 Plot between Distance and Time

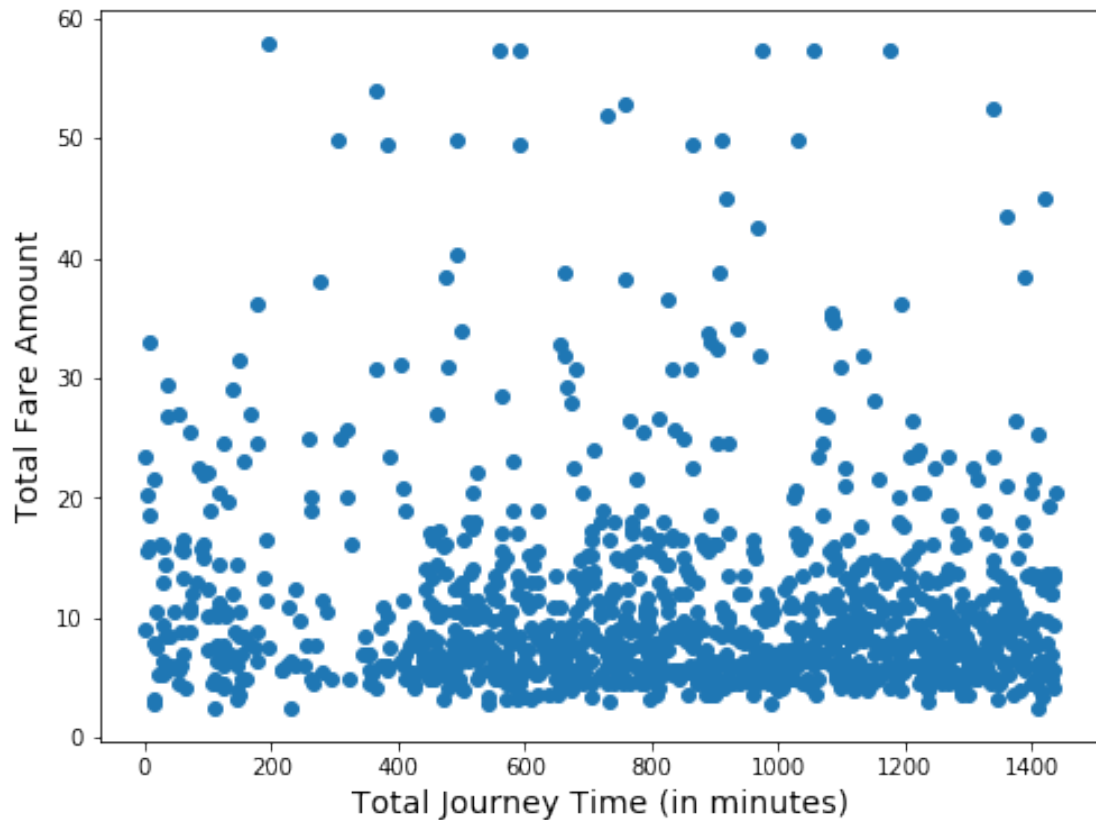
```
In [19]: plt.figure(figsize = (8,6))
plt.xlabel('Distance Traveled', fontsize = 14)
plt.ylabel('Total Journey Time (in minutes)', fontsize = 14)
plt.scatter(train_df[:1200].distance, train_df[:1200].journey_time)
plt.show()
```



- The plot between Distance traveled and the total journey time of a passenger generates a non-linear relationship. This plot doesn't tell us the exact relationship between the variables

1.0.6 Plot between Total Fare Amount and Time

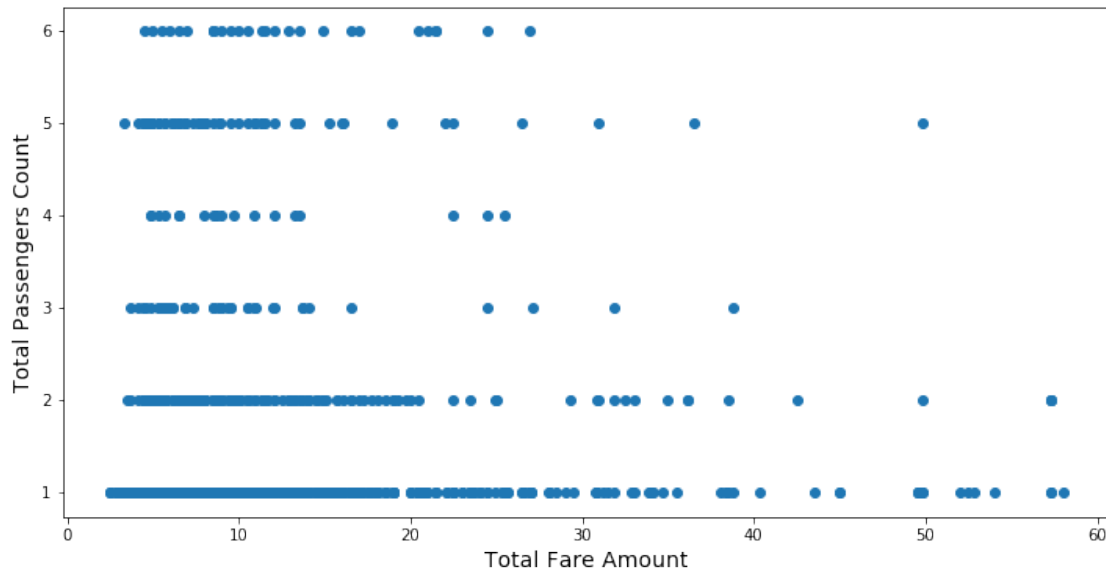
```
In [20]: plt.figure(figsize = (8,6))
plt.xlabel('Total Journey Time (in minutes)', fontsize = 14)
plt.ylabel('Total Fare Amount', fontsize = 14)
plt.scatter(train_df[:1200].journey_time, train_df[:1200].fare_amount)
plt.show()
```



- The plot between Distance traveled and the total journey time of a passenger generates a non - linear relationship. The plot gives a scattered data and hence we cannot infer anything from it

1.0.7 Plot between Total Fare Amount and Total Passenger Count

```
In [21]: plt.figure(figsize = (12,6))
plt.xlabel('Total Fare Amount', fontsize = 14)
plt.ylabel('Total Passengers Count', fontsize = 14)
plt.scatter(train_df[:1200].fare_amount, train_df[:1200].passenger_count)
plt.show()
```




- The plot between total fare and the number of passengers taking the cab, we see that generally people in New York City are spending less than 30\$ while they take the cab service, although it is a very generic statement

1.0.8 Bar Chart comparing the fare amount depending on the hour of the day

```
In [22]: pickup_hour_time = (train_df['pickup_datetime'].dt.hour)[:1200]
         fare = train_df[:1200].fare_amount

plt.figure(figsize = (10,6))
plt.bar(pickup_hour_time, fare)
plt.xlabel('Hour of the day')
plt.ylabel('Fare amount')
plt.show()
```




- This plot shows relationship between the fare amount at particular hours of the day. We can see at what times during the day, we have highest amount of fare. It seems that during early morning hours(3 AM, 9 AM), and the evening at 5PM, 6PM and 8 PM we are able to see the highest fare amount throughout the day.

1.0.9 Data plot based on Pickup locations

```
In [23]: plt.figure(figsize = (12,8))
plt.title('Pickup Locations', fontsize=14)
p_long, p_lat = pd.Series(train_df[:1000000].pickup_longitude, name="Pickup Longitude")

pickup = sb.regplot(x=p_long, y=p_lat, scatter_kws={'alpha':0.3})
pickup.set(xlim = (-74.1, -73.7))
```


```
Out[23]: [(-74.1, -73.7)]
```

- From the given plot we get to know the various locations of pickups based on their latitude and longitude coordinates. Since we have a huge dataset, we can see the whole map of New York City from the marking locations of pickup. We can also see there are a few points for the pickup in the right bottom corner, because those coordinates are of JFK airport, while a few scattered points are also present which represent the pickup locations in the New York's boroughs as well. So we can take the pickup based on these locations as well.

1.0.10 Data plot based on DropOff locations

```
In [24]: plt.figure(figsize = (12,8))
plt.title('Dropoff Locations', fontsize=14)
d_long, d_lat = pd.Series(train_df[:1000000].dropoff_longitude, name="Dropoff Longitude")
dropoff = sb.regplot(x=d_long, y=d_lat, marker="+")
```



- From the given plot we get to know the various locations of dropoff points based on their latitude and longitude coordinates. We can see the map of New York City from the marking locations of dropoff. We can also see there are a few points for the pickup in the right bottom corner, because those coordinates are of JFK airport. Here we see a lot of scattered points across Manhattan, so we can say that, many of the taxi cab's customers have a drop off location outside of New York City.

1.0.11 Additional Feature Extraction

```
In [25]: train_df['day'] = train_df['pickup_datetime'].dt.day;
train_df['month'] = train_df['pickup_datetime'].dt.month;
train_df['hour'] = train_df['pickup_datetime'].dt.hour;
train_df['minute'] = train_df['pickup_datetime'].dt.minute;
train_df.head()
```

```
Out[25]:
```

	key	fare_amount	pickup_datetime \
0	2009-06-15 17:26:21.0000001	4.5	2009-06-15 17:26:21
2	2011-08-18 00:35:00.00000049	5.7	2011-08-18 00:35:00
3	2012-04-21 04:30:42.0000001	7.7	2012-04-21 04:30:42
4	2010-03-09 07:51:00.000000135	5.3	2010-03-09 07:51:00
5	2011-01-06 09:50:45.0000002	12.1	2011-01-06 09:50:45

	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude \
--	------------------	-----------------	-------------------	--------------------

0	-73.844311	40.721319	-73.841610	40.712278
2	-73.982738	40.761270	-73.991242	40.750562
3	-73.987130	40.733143	-73.991567	40.758092
4	-73.968095	40.768008	-73.956655	40.783762
5	-74.000964	40.731630	-73.972892	40.758233

	passenger_count	distance	journey_time	day	month	hour	minute
0	1	0.009436	1046	15	6	17	26
2	2	0.013674	35	18	8	0	35
3	1	0.025340	270	21	4	4	30
4	1	0.019470	471	9	3	7	51
5	1	0.038675	590	6	1	9	50

1.0.12 Training Data - Linear Regression

```
In [38]: data2 = train_df[:5000000]
```

```
#We are considering three features for the data selection. These three features could be
features = ['passenger_count', 'distance', 'journey_time']
X = data2[features]
y = data2['fare_amount']
```

```
#Source for Linear Regression: https://towardsdatascience.com/linear-regression-in-python
from sklearn.cross_validation import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
```

```
#We are splitting the data for training and testing according to the ratio of 80 : 20
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
```

```
lr = LinearRegression()
lr.fit(X_train, y_train)
```

```
(4000000, 3) (1000000, 3) (4000000,) (1000000,)
```

```
Out[38]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

```
In [39]: print(lr.intercept_)
          print(lr.coef_)
          zip(features, lr.coef_)
```

```
3.680830954073837
```

```
[3.75550574e-02 2.14014535e+02 2.38247046e-04]
```

```
Out[39]: <zip at 0x7f56051f1ac8>
```

```
In [40]: y_pred = lr.predict(X_test)
         print(y_pred)
         print("RMS: ", np.sqrt(metrics.mean_squared_error(y_test, y_pred)))

[10.47004978 20.54124327  6.02887277 ... 11.37778047  4.58511403
  7.04375718]
RMS:  3.91867263627871
```

We are getting RMS value = 3.912 from the training data set

1.0.13 Actual Data Prediction using Linear Regression

```
In [41]: actual_data = pd.read_csv('./test.csv')
         actual_data.dtypes
```

```
Out[41]: key                object
         pickup_datetime    object
         pickup_longitude  float64
         pickup_latitude   float64
         dropoff_longitude  float64
         dropoff_latitude   float64
         passenger_count    int64
         dtype: object
```

```
In [42]: eu_cal = (actual_data['dropoff_latitude'] - actual_data['pickup_latitude']) **2 + (act
         eu_dist = np.sqrt(eu_cal)
```

```
In [43]: actual_data['pickup_datetime'] = actual_data['pickup_datetime'].str.replace(" UTC", "")
         #replace the given date time in a new format
         actual_data['pickup_datetime'] = pd.to_datetime(actual_data['pickup_datetime'], format=
```

```
In [44]: journey_time = (actual_data['pickup_datetime'].dt.hour)*60 + actual_data['pickup_dateti
```

```
In [45]: actual_data['distance'] = eu_dist
         actual_data['journey_time'] = journey_time
```

```
In [46]: actual_data['day'] = actual_data['pickup_datetime'].dt.day;
         actual_data['month'] = actual_data['pickup_datetime'].dt.month;
         actual_data['hour'] = actual_data['pickup_datetime'].dt.hour;
         actual_data['minute'] = actual_data['pickup_datetime'].dt.minute;
         actual_data.head()
```

```
Out[46]:
```

	key	pickup_datetime	pickup_longitude	\
0	2015-01-27 13:08:24.0000002	2015-01-27 13:08:24	-73.973320	
1	2015-01-27 13:08:24.0000003	2015-01-27 13:08:24	-73.986862	
2	2011-10-08 11:53:44.0000002	2011-10-08 11:53:44	-73.982524	
3	2012-12-01 21:12:12.0000002	2012-12-01 21:12:12	-73.981160	

```

4  2012-12-01 21:12:12.0000003 2012-12-01 21:12:12          -73.966046

    pickup_latitude  dropoff_longitude  dropoff_latitude  passenger_count  \
0      40.763805      -73.981430      40.743835          1
1      40.719383      -73.998886      40.739201          1
2      40.751260      -73.979654      40.746139          1
3      40.767807      -73.990448      40.751635          1
4      40.789775      -73.988565      40.744427          1

    distance  journey_time  day  month  hour  minute
0  0.021554      788      27      1     13      8
1  0.023180      788      27      1     13      8
2  0.005870      713      8      10     11     53
3  0.018649     1272      1      12     21     12
4  0.050631     1272      1      12     21     12

```

```

In [47]: features1 = ['passenger_count', 'distance', 'journey_time']
        X1 = actual_data[features1]

```

```

In [49]: predict_value = lr.predict(X1)
        print(predict_value)

```

```

[ 8.51897634  8.86693473  5.14460608 ... 50.6405486  21.23970975
 6.95370272]

```

```

In [50]: final_data = pd.DataFrame()
        final_data['key'] = actual_data['key']
        final_data['fare_amount'] = predict_value
        final_data.to_csv('final_result.csv', sep=',', index = False)

```

1.0.14 Training Data - Random Forest Regressor

```

In [167]: data3 = train_df[:5000000]

```

```

from sklearn.ensemble import RandomForestRegressor
from sklearn.datasets import make_regression

```

```

In [168]: regressor = RandomForestRegressor(max_depth= 12, random_state=0,n_estimators=5)

```

```

#We are considering three features for the data selection. These three features could
#I considered other features as well, but they didn't improve the score at all. These
features3 = ['passenger_count', 'distance', 'journey_time']

```

```

X3 = data3[features3]
y3 = data3['fare_amount']

```

```

regressor.fit(X3, y3)

```

```

#Source: http://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForests

```

```
Out[168]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=12,
                                max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=5, n_jobs=1,
                                oob_score=False, random_state=0, verbose=0, warm_start=False)
```

```
In [169]: print(regressor.feature_importances_)
```

```
[6.96457525e-04 9.87749333e-01 1.15542096e-02]
```

```
In [170]: #We are taking the testing data from the current given data set - for local prediction
X3_test = train_df[5000001: 6000000][features3]
```

```
predict3 = regressor.predict(X3_test)
print(predict3)
```

```
[ 5.67247324  5.78835029  5.36457978 ...  6.13259908 11.51195472
  5.36457978]
```

1.0.15 Data Prediction using Random Forest

```
In [171]: actual_data3 = pd.read_csv('./test.csv')
actual_data3.dtypes
```

```
Out[171]: key                object
pickup_datetime            object
pickup_longitude           float64
pickup_latitude            float64
dropoff_longitude          float64
dropoff_latitude           float64
passenger_count            int64
dtype: object
```

```
In [172]: eu_cal3 = (actual_data3['dropoff_latitude'] - actual_data3['pickup_latitude']) **2 +

eu_dist3 = np.sqrt(eu_cal3)
```

```
In [173]: actual_data3['pickup_datetime'] = actual_data3['pickup_datetime'].str.replace(" UTC",
#replace the given date time in a new format
actual_data3['pickup_datetime'] = pd.to_datetime(actual_data3['pickup_datetime'], format=
```

```
In [174]: #Calculating the time using hour and minute of the given time stamp
journey_time3 = (actual_data3['pickup_datetime'].dt.hour)*60 + actual_data3['pickup_da
```

```
In [175]: actual_data3['distance'] = eu_dist3
actual_data3['journey_time'] = journey_time3
```

```
In [176]: #Updating the data and adding new columns to accomodate new features
actual_data3['day'] = actual_data3['pickup_datetime'].dt.day;
actual_data3['month'] = actual_data3['pickup_datetime'].dt.month;
actual_data3['hour'] = actual_data3['pickup_datetime'].dt.hour;
actual_data3['minute'] = actual_data3['pickup_datetime'].dt.minute;
actual_data3.head()
```

```
Out[176]:
```

	key	pickup_datetime	pickup_longitude	\
0	2015-01-27 13:08:24.0000002	2015-01-27 13:08:24	-73.973320	
1	2015-01-27 13:08:24.0000003	2015-01-27 13:08:24	-73.986862	
2	2011-10-08 11:53:44.0000002	2011-10-08 11:53:44	-73.982524	
3	2012-12-01 21:12:12.0000002	2012-12-01 21:12:12	-73.981160	
4	2012-12-01 21:12:12.0000003	2012-12-01 21:12:12	-73.966046	

	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	\
0	40.763805	-73.981430	40.743835	1	
1	40.719383	-73.998886	40.739201	1	
2	40.751260	-73.979654	40.746139	1	
3	40.767807	-73.990448	40.751635	1	
4	40.789775	-73.988565	40.744427	1	

	distance	journey_time	day	month	hour	minute
0	0.021554	788	27	1	13	8
1	0.023180	788	27	1	13	8
2	0.005870	713	8	10	11	53
3	0.018649	1272	1	12	21	12
4	0.050631	1272	1	12	21	12

```
In [177]: X31 = actual_data3[features3]
predict_value3 = regressor.predict(X31)
```

```
In [178]: final_data3 = pd.DataFrame()
final_data3['key'] = actual_data3['key']
final_data3['fare_amount'] = predict_value3
final_data3.to_csv('final_result3.csv',sep=',', index = False)
```